



**Rayat Shikshan Sanstha's
KARMAVEER BHAURAO PATIL COLLEGE, VASHI
(Empowered Autonomous)**



Reaccredited NAAC with Grade 'A++' (CGPA3.51)|ISO 9001:2015 Certified Institute
'Best College' Award by University of Mumbai

[DEPARTMENT OF INFORMATION TECHNOLOGY]

CERTIFICATE

This is to certify that **Mr. SHUBHAM VIKAS NAVALE**, student of Class **MScIT – II** from **Karmaveer Bhaurao Patil College, Vashi, Navi Mumbai** has satisfactorily completed the practical course in subject **Machine Learning**, as per the syllabus laid by the college during the Academic Year **2025-26**.

ROLL NO.: 25176019

EXAM NO.:25176019

Date:12/11/2025

MANOJ CHOUDHARY

Course Coordinator

MADHURI GABHANE

Head of Department
Information Technology

External Examiner

Index

Sr No	Name of Practicals	Date	Signature
1	Diabetes Prediction Using Pima Indians Dataset		
2	Student Academic Performance Prediction (Pass/Fail Classification)		
3	Customer Churn Prediction for Telecom Industry		
4	Email Spam Detection and Classification		
5	Time-Series Anomaly Detection Techniques		
6	Human Activity Recognition Using Sensor Data		
7	Wine Quality Prediction Using Machine Learning		
8	Loan Eligibility Prediction Model		

Practical No :- 1

Aim:- Pima Indians Diabetes Prediction

Code:-

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, confusion_matrix,
classification_report

# --- Load dataset ---
url = "https://raw.githubusercontent.com/jbrownlee/Datasets/master/pima-
indians-diabetes.data.csv"
columns = [
    'Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness',
    'Insulin', 'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'
]
data = pd.read_csv(url, header=None, names=columns)

# --- Explore dataset ---
print("First 5 rows:\n", data.head())
print("\nInfo:\n")
print(data.info())
print("\nSummary Statistics:\n", data.describe())

# --- Clean data: replace zeros with median ---
```

```
cols_to_fix = ['Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI']

for col in cols_to_fix:
    data[col] = data[col].replace(0, np.nan)
    data[col] = data[col].fillna(data[col].median())

# --- Correlation heatmap ---
plt.figure(figsize=(10, 8))
sns.heatmap(data.corr(numeric_only=True), annot=True, cmap='coolwarm')
plt.title("Correlation Matrix")
plt.tight_layout()
plt.show()

# --- Outcome distribution ---
plt.figure(figsize=(6, 4))
sns.countplot(x='Outcome', data=data)
plt.title("Outcome Distribution")
plt.tight_layout()
plt.show()

# --- Prepare training data ---
X = data.drop('Outcome', axis=1)
y = data['Outcome']
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

# --- Train model ---
```

```
model = RandomForestClassifier(random_state=42)
model.fit(X_train, y_train)
y_pred = model.predict(X_test)

# --- Evaluate model ---
accuracy = accuracy_score(y_test, y_pred)
print(f"\nModel Accuracy: {accuracy * 100:.2f}%")
print("\nConfusion Matrix:\n", confusion_matrix(y_test, y_pred))
print("\nClassification Report:\n", classification_report(y_test, y_pred))

# --- Feature importance plot ---
importances = model.feature_importances_
feat_importances = pd.Series(importances, index=X.columns)
plt.figure(figsize=(10, 6))
feat_importances.sort_values().plot(kind='barh')
plt.title("Feature Importance from Random Forest")
plt.tight_layout()
plt.show()

# --- Predict for new sample patient ---
new_patient_data = pd.DataFrame([
    'Pregnancies': 2,
    'Glucose': 120,
    'BloodPressure': 70,
    'SkinThickness': 30,
    'Insulin': 80,
    'BMI': 25.6,
```

```

'DiabetesPedigreeFunction': 0.6,
'Age': 32
})]

prediction = model.predict(new_patient_data)
print("\nPrediction for new patient:", "Likely diabetic." if prediction[0] == 1
else "Unlikely diabetic.")Output:-
```

```

First 5 rows:
   Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin  BMI \
0           6      148             72            35       0  33.6
1           1       85             66            29       0  26.6
2           8      183             64            0       0  23.3
3           1       89             66            23      94  28.1
4           0      137             40            35     168  43.1

   DiabetesPedigreeFunction  Age  Outcome
0           0.627    50       1
1           0.351    31       0
2           0.672    32       1
3           0.167    21       0
4           2.288    33       1

Info:


<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Pregnancies      768 non-null    int64  
 1   Glucose          768 non-null    int64  
 2   BloodPressure    768 non-null    int64  
 3   SkinThickness    768 non-null    int64  
 4   Insulin          768 non-null    int64  
 5   BMI              768 non-null    float64 
 6   DiabetesPedigreeFunction 768 non-null    float64 
 7   Age              768 non-null    int64  
 8   Outcome          768 non-null    int64  
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
None
```



```

Summary Statistics:
   Pregnancies      Glucose  BloodPressure  SkinThickness      Insulin \
count    768.000000  768.000000  768.000000  768.000000  768.000000
mean     3.845052  120.894531   69.105469   20.536458  79.799479
std      3.369578   31.972618   19.355807   15.952218 115.244002
min      0.000000   0.000000   0.000000   0.000000   0.000000
25%     1.000000   99.000000   62.000000   0.000000   0.000000
50%     3.000000  117.000000   72.000000  23.000000  30.500000
75%     6.000000  140.250000   80.000000  32.000000 127.250000
max    17.000000  199.000000  122.000000  99.000000 846.000000

      BMI  DiabetesPedigreeFunction      Age      Outcome
count  768.000000  768.000000  768.000000  768.000000
mean   31.992578   0.471876  33.240885   0.348958
std    7.884160    0.331329  11.760232   0.476951
min    0.000000    0.078000  21.000000   0.000000
25%   27.300000   0.243750  24.000000   0.000000
50%   32.000000   0.372500  29.000000   0.000000
75%   36.600000   0.626250  41.000000   1.000000
max   67.100000   2.420000  81.000000   1.000000
/tmpp/ipython-input-966236812.py:29: FutureWarning: A value is trying to be set on a copy.
The behavior will change in pandas 3.0. This inplace method will never work because the :

```





Practical No :- 2

Aim:- Student Academic Performance (Pass/Fail Prediction)

Code:-

```
import pandas as pd  
from sklearn.model_selection import train_test_split  
from sklearn.preprocessing import MinMaxScaler  
from sklearn.ensemble import RandomForestClassifier  
from sklearn.metrics import accuracy_score, classification_report  
  
# Load data  
df = pd.read_csv('StudentsPerformance.csv')  
  
# Drop 'StudentID' if it exists  
df = df.drop(columns=['StudentID'], errors='ignore')  
  
# Create 'Pass_Fail' based on average score  
passing_marks = 60  
df['Pass_Fail'] = (df[['math score', 'reading score', 'writing score']].mean(axis=1)  
>= passing_marks).astype(int)  
  
# Encode categorical columns  
df = pd.get_dummies(df, drop_first=True)  
  
# Features and target  
X = df.drop('Pass_Fail', axis=1)  
y = df['Pass_Fail']
```

```

# Scale features
scaler = MinMaxScaler()
X_scaled = scaler.fit_transform(X)

# Train/test split
X_train, X_test, y_train, y_test = train_test_split(
    X_scaled, y, test_size=0.2, random_state=42
)

# Train model
model = RandomForestClassifier(random_state=42)
model.fit(X_train, y_train)

# Evaluation
y_pred = model.predict(X_test)
print("Accuracy:", accuracy_score(y_test, y_pred))
print(classification_report(y_test, y_pred))

```

Output:-

Accuracy: 0.98				
	precision	recall	f1-score	support
0	0.95	0.98	0.97	62
1	0.99	0.98	0.99	138
accuracy			0.98	200
macro avg			0.97	200
weighted avg			0.98	200

Practical No:- 3

Aim:- Customer Churn Prediction (Telco)

Code:-

```
import pandas as pd

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix

# Load
df = pd.read_csv('P3WA_Fn-UseC_-Telco-Customer-Churn.csv')

# Clean
df['TotalCharges'] = pd.to_numeric(df['TotalCharges'], errors='coerce')
# Fix: Avoid chained assignment warning by assigning back the filled column
df['TotalCharges'] = df['TotalCharges'].fillna(df['TotalCharges'].median())
df.drop('customerID', axis=1, inplace=True)

# Encode
df['Churn'] = df['Churn'].map({'Yes':1, 'No':0})
df.replace({'No phone service':'No', 'No internet service':'No'}, inplace=True)
df = pd.get_dummies(df, drop_first=True)

# Split
X = df.drop('Churn', axis=1)
y = df['Churn']
```

```
X_train, X_test, y_train, y_test = train_test_split(  
    X, y, test_size=0.2, random_state=42, stratify=y  
)  
  
# Scale  
num_cols = ['tenure', 'MonthlyCharges', 'TotalCharges']  
scaler = StandardScaler()  
X_train[num_cols] = scaler.fit_transform(X_train[num_cols])  
X_test[num_cols] = scaler.transform(X_test[num_cols])  
  
# Model  
clf = RandomForestClassifier(random_state=42)  
clf.fit(X_train, y_train)  
y_pred = clf.predict(X_test)  
  
# Evaluate  
print("Accuracy:", accuracy_score(y_test, y_pred))  
print(classification_report(y_test, y_pred))  
print(confusion_matrix(y_test, y_pred))  
Output:-
```

```
ipython-stdin-1682618412.py:12: FutureWarning: A value is trying to be set
The behavior will change in pandas 3.0. This inplace method will never work because
For example, when doing 'df[col].method(value, inplace=True)', try using 'df.met
    df['TotalCharges'].fillna(df['TotalCharges'].median(), inplace=True)
Accuracy: 0.794180269694819
      precision    recall  f1-score   support
          0       0.83     0.90     0.87     1035
          1       0.65     0.49     0.56      374
    accuracy         0.79
  macro avg       0.74     0.70     0.71     1409
weighted avg       0.78     0.79     0.78     1409
[[934 101]
 [189 185]]
```

Practical No :- 4

Aim:- Email Spam Classifier

Code:-

```
import pandas as pd  
from sklearn.model_selection import train_test_split  
from sklearn.feature_extraction.text import TfidfVectorizer  
from sklearn.naive_bayes import MultinomialNB  
from sklearn.metrics import accuracy_score, classification_report  
  
# Load (example: SMS Spam CSV)  
df = pd.read_csv('spam.csv', encoding='latin-1')[['v1','v2']]  
df.columns = ['label', 'text']  
df['label'] = df['label'].map({'spam':1, 'ham':0})  
  
X_train, X_test, y_train, y_test = train_test_split(  
    df['text'], df['label'], test_size=0.2, random_state=42  
)  
  
vectorizer = TfidfVectorizer(stop_words='english')  
X_train_tfidf = vectorizer.fit_transform(X_train)  
X_test_tfidf = vectorizer.transform(X_test)  
  
clf = MultinomialNB()  
clf.fit(X_train_tfidf, y_train)  
y_pred = clf.predict(X_test_tfidf)  
  
print("Accuracy:", accuracy_score(y_test, y_pred))
```

```
print(classification_report(y_test, y_pred))
```

Output:-

Accuracy: 0.9668161434977578					
	precision	recall	f1-score	support	
0	0.96	1.00	0.98	965	
1	1.00	0.75	0.86	150	
accuracy			0.97	1115	
macro avg	0.98	0.88	0.92	1115	
weighted avg	0.97	0.97	0.96	1115	

Practical No :- 5

Aim:- Time-Series Anomaly Detection

Code:-

```
import pandas as pd
import numpy as np
from sklearn.ensemble import IsolationForest

df = pd.read_csv('your_timeseries.csv', parse_dates=['timestamp'],
index_col='timestamp')
series = df['value']

model = IsolationForest(contamination=0.01, random_state=42)
df['anomaly'] = model.fit_predict(series.values.reshape(-1,1))
anomalies = df[df['anomaly'] == -1]
print("Anomalies found:", anomalies)
```

Output:-

Anomalies found:		value	anomaly
timestamp			
2025-01-01 09:00:00	1000		-1

Practical No :- 6

Aim:- Human Activity Recognition (HAR)

Code:-

```
import pandas as pd  
import numpy as np  
from sklearn.model_selection import train_test_split  
from sklearn.ensemble import RandomForestClassifier  
from sklearn.metrics import classification_report
```

```
# --- Example dataset creation ---  
# 100 samples, 6 sensor features  
np.random.seed(42)  
X = pd.DataFrame({  
    'accel_x': np.random.randn(100),  
    'accel_y': np.random.randn(100),  
    'accel_z': np.random.randn(100),  
    'gyro_x': np.random.randn(100),  
    'gyro_y': np.random.randn(100),  
    'gyro_z': np.random.randn(100),  
})
```

```
# Example activity labels  
activities = ['Walking', 'Running', 'Sitting']  
y = pd.DataFrame({  
    'Activity': np.random.choice(activities, size=100)  
})
```

```
# Save to CSV (optional)
X.to_csv('X.csv', index=False)
y.to_csv('y.csv', index=False)
print("□ Files 'X.csv' and 'y.csv' created successfully.")

# --- Load CSVs (simulating real dataset) ---
X = pd.read_csv('X.csv')
y = pd.read_csv('y.csv')['Activity']

# Train/test split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

# Train Random Forest classifier
clf = RandomForestClassifier(random_state=42)
clf.fit(X_train, y_train)

# Predict and evaluate
y_pred = clf.predict(X_test)
print("\nClassification Report:\n")
print(classification_report(y_test, y_pred))
```

Output:-

→ ✓ Files 'X.csv' and 'y.csv' created successfully.

Classification Report:

	precision	recall	f1-score	support
Running	0.50	0.29	0.36	7
Sitting	0.71	0.71	0.71	7
Walking	0.33	0.50	0.40	6
accuracy			0.50	20
macro avg	0.52	0.50	0.49	20
weighted avg	0.53	0.50	0.50	20

Practical N0:- 7

Aim:- Wine Quality Prediction

Code:-

```
import pandas as pd  
from sklearn.model_selection import train_test_split  
from sklearn.ensemble import RandomForestClassifier  
from sklearn.metrics import accuracy_score, classification_report  
  
df = pd.read_csv('winequality-red.csv', sep=';')  
df['good_quality'] = (df['quality'] >= 6).astype(int)  
  
X = df.drop(['quality', 'good_quality'], axis=1)  
y = df['good_quality']  
  
X_train, X_test, y_train, y_test = train_test_split(  
    X, y, test_size=0.2, random_state=42  
)  
  
clf = RandomForestClassifier(random_state=42)  
clf.fit(X_train, y_train)  
y_pred = clf.predict(X_test)  
  
print("Accuracy:", accuracy_score(y_test, y_pred))  
print(classification_report(y_test, y_pred))
```

Output:-

```
→ Accuracy: 0.5
      precision    recall   f1-score   support
          0       0.50      1.00      0.67       1
          1       0.00      0.00      0.00       1

   accuracy                           0.50      2
macro avg       0.25      0.50      0.33      2
weighted avg    0.25      0.50      0.33      2

/usr/local/lib/python3.12/dist-packages/sklearn/metrics/
    _warn_prf(average, modifier, f"{metric.capitalize()} is "
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/
    _warn_prf(average, modifier, f"{metric.capitalize()} is "
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/
    _warn_prf(average, modifier, f"{metric.capitalize()} is "
```

Practical No:- 8

Aim:- Loan Eligibility Prediction

Code:-

```
import pandas as pd  
from sklearn.model_selection import train_test_split  
from sklearn.ensemble import RandomForestClassifier  
from sklearn.metrics import accuracy_score, classification_report  
  
# Load and preprocess data  
df = pd.read_csv('loan_data.csv')  
df['Loan_Status'] = df['Loan_Status'].map({'Y': 1, 'N': 0})  
df.drop(['Loan_ID'], axis=1, inplace=True)  
df = pd.get_dummies(df, drop_first=True)  
  
X = df.drop('Loan_Status', axis=1)  
y = df['Loan_Status']  
  
# Stratified split with larger test size  
X_train, X_test, y_train, y_test = train_test_split(  
    X, y, test_size=0.3, random_state=42, stratify=y  
)  
  
# Print class distributions  
print("Train class distribution:\n", y_train.value_counts())  
print("Test class distribution:\n", y_test.value_counts())  
  
# Train model
```

```

clf = RandomForestClassifier(random_state=42)
clf.fit(X_train, y_train)


# Predict on test set
y_pred = clf.predict(X_test)

# Print actual vs predicted labels
print("Actual labels:", y_test.values)
print("Predicted labels:", y_pred)

# Print accuracy and classification report (with zero_division=0 to avoid
warnings)
print("Accuracy:", accuracy_score(y_test, y_pred))
print(classification_report(y_test, y_pred, zero_division=0))

```

Output:-


A screenshot of a terminal window displaying a classification report. The report includes metrics for two classes (0 and 1) and overall averages. The 'accuracy' row shows 0.50. The 'macro avg' row shows precision, recall, f1-score, and support all at 0.25. The 'weighted avg' row shows precision, recall, f1-score, and support all at 0.33. The report is preceded by a warning message about deprecated code related to 'warn_prf' and 'metric.capitalize()'.